



Terry Dhariwal

Your cost per job
report is **missing**
the jobs that cost
the most.

5 steps to find them

(straight from Databricks' own docs)



Why most cost per job reports leave the big ones out

Databricks attributes cost to a job using the **job ID** in billing data.

Jobs that run on **All-Purpose clusters** don't carry one.

So Databricks' own cost-per-job queries exclude them. And All-Purpose is the priciest classic compute for a scheduled job.



1. Spot the gap

In **system.billing.usage**, the job ID is only filled in for:

- + *Jobs on job compute*
- + *Serverless jobs*
- × *Jobs on All-Purpose clusters*

If a big share of your spend has no job ID, this is usually why.



2. Price the gap

Same work, different rate. AWS
Premium list price per DBU:

\$0.55

All-Purpose
Classic

\$0.15

Jobs
Classic

That's **3.67x the DBU rate** for the
same job. The cloud VMs cost the
same either way.



3. Find the jobs

- 1 *List every scheduled job and its tasks*
- 2 *Flag tasks set to `existing_cluster_id`*
- 3 *Those run on All-Purpose compute*
- 4 *Note each cluster's owner and tags*

Most teams find a handful of them.
They're usually old, busy and
nobody's favourite.



4. Move them

Point scheduled work at **job compute** or **serverless jobs**.

- + *Every run gets a job ID*
- + *Cost per job becomes real*
- + *The rate per DBU drops*

It's configuration, not a migration. Job clusters add start-up time, so use pools or serverless where that matters.



5. Report the rest honestly

Some spend will still have no job: notebooks, shared clusters, ad hoc work.

Report it as a **visibility gap**. Give it an owner and a target.

Just never count it as savings.
Untraced isn't the same as wasted.



BigHammer.ai

Want to see this on a real bill?

**Reduce your
Databricks costs
by up to 75%**

Free live masterclass
Thu 15 Oct · 12:00 PM ET · 5 PM UK

Hosted by **Srinath Reddy**,
founder and CEO of BigHammer.ai



webinar.bighammerai.com